


C Boot Camp

Oct. 1, 2017

Aarohi
Fernando
Jiayi
Satoru
Stephen



Agenda

- C Basics
- Debugging Tools / Demo
- Appendix
 - C Standard Library
 - getopt
 - stdio.h
 - stdlib.h
 - string.h



C Basics Handout

```
ssh <andrewid>@shark.ics.cs.cmu.edu  
cd ~/private  
wget http://cs.cmu.edu/~213/activities/cbootcamp.tar.gz  
tar xvpf cbootcamp.tar.gz  
cd cbootcamp  
make
```

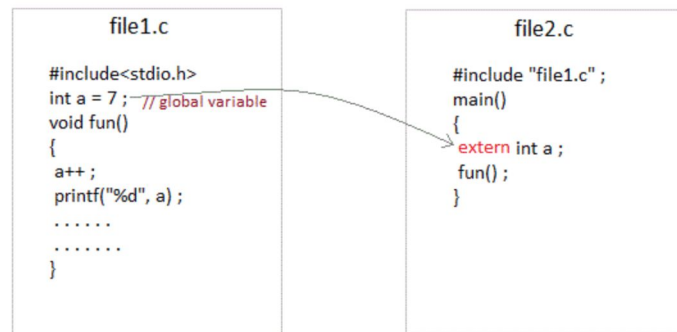
- Contains useful, self-contained C examples
- Slides relating to these examples will have the file names in the **top-right corner!**

C Basics

- The *minimum* you must know to do well in this class
 - You have seen these concepts before
 - Make sure you remember them.
- Summary:
 - Pointers/Arrays/Structs/Casting
 - Memory Management
 - Function pointers/Generic Types
 - Strings

Variable Declarations & Qualifiers

- **Global Variables:**
 - Defined outside functions, seen by all files
 - Use “extern” keyword to use a global variable defined in another file
- **Const Variables:**
 - For variables that won't change
 - Data stored in read-only data section
- **Static Variables:**
 - For locals, keeps value between invocations
 - USE SPARINGLY
 - Note: static has a different meaning when referring to functions (not visible outside of object file)



global variable from one file can be used in other using **extern** keyword.

```

#include<stdio.h>
int fun()
{
    static int count = 0;
    count++;
    return count;
}

int main()
{
    printf("%d ", fun());
    printf("%d ", fun());
    return 0;
}

```

Output:

Casting

- Can convert a variable to a different type
- Integer Casting:
 - Signed <-> Unsigned: Keep Bits - Re-Interpret
 - Small -> Large: Sign-Extend MSB, preserve value
- Cautions:
 - Cast Explicitly: `int x = (int) y` instead of `int x = y`
 - Casting Down: Truncates data
 - Cast Up: Upcasting and dereferencing a pointer causes undefined memory access
- Rules for Casting Between Integer Types


Pointers

- Stores address of a value in memory
 - e.g. `int*`, `char*`, `int**`, etc
 - Access the value by dereferencing (e.g. `*a`).
Can be used to read or write a value to given address
 - Dereferencing `NULL` causes undefined behavior (usually a segfault)



Pointers

- Pointer to type A references a block of `sizeof(A)` bytes
- Get the address of a value in memory with the `&` operator
- Pointers can be *aliased*, or pointed to same address



Pointer Arithmetic

./pointers

- Can add/subtract from an address to get a new address
 - Only perform when absolutely necessary (i.e., malloc/strdup)
 - Result depends on the pointer type
- $A+i$, where A is a pointer = 0x100, i is an int
 - $\text{int}^* A: A+i = 0x100 + \text{sizeof}(\text{int}) * i = 0x100 + 4 * i$
 - $\text{char}^* A: A+i = 0x100 + \text{sizeof}(\text{char}) * i = 0x100 + 1 * i$
 - $\text{int}^{**} A: A+i = 0x100 + \text{sizeof}(\text{int}^*) * i = 0x100 + 8 * i$
- Rule of thumb: **explicitly** cast pointer to avoid confusion
 - Prefer $((\text{char}^*) (A) + i)$ to $(A + i)$, even if A has type char^*

Pointer Arithmetic

`./pointers`

- The 'pointers' program demonstrates how values of different sizes can be written to and read back from the memory.
- The examples are to show you how the ~type~ of the pointer affects arithmetic done on the pointer.
- When adding x to a pointer A (i.e. $A + x$), the result is really $(A + x * \text{sizeof}(\text{TYPE_OF_PTR_A}))$.
- Run the 'pointers' program
`$./pointers`

Call by Value vs Call by Reference

- Call-by-value: Changes made to arguments passed to a function *aren't* reflected in the calling function
- Call-by-reference: Changes made to arguments passed to a function *are* reflected in the calling function
- C is a **call-by-value** language
- To cause changes to values outside the function, use pointers
 - Do *not* assign the pointer to a different value (that won't be reflected!)
 - Instead, *dereference the pointer* and assign a value to that address


```
void swap(int* a, int* b) {  
    int temp = *a;  
    *a = *b;  
    *b = temp;  
}
```

```
int x = 42;  
int y = 54;  
swap(&x, &y);  
printf("%d\n", x); // 54  
printf("%d\n", y); // 42
```

Arrays/Strings

- Arrays: fixed-size collection of elements of the same type
 - Can allocate on the stack or on the heap
 - `int A[10]; // A is array of 10 int's on the stack`
 - `int* A = calloc(10, sizeof(int)); // A is array of 10 int's on the heap`
- Strings: Null-character ('\0') terminated character arrays
 - Null-character tells us where the string ends
 - All standard C library functions on strings assume null-termination.

C Program Memory Layout



Stack vs Heap vs Data

- Local variables and function arguments are placed on the *stack*
 - deallocated after the variable leaves scope
 - *do not* return a pointer to a stack-allocated variable!
 - *do not* reference the address of a variable outside its scope!
- Memory blocks allocated by calls to malloc/calloc are placed on the *heap*
- Globals, constants are placed in *data* section
- Example:
 - `// a is a pointer on the stack to a memory block on the heap`
 - `int* a = malloc(sizeof(int));`

Malloc, Free, Calloc

- Handle dynamic memory allocation on HEAP
- `void* malloc (size_t size):`
 - allocate block of memory of `size` bytes
 - does not initialize memory
- `void* calloc (size_t num, size_t size):`
 - allocate block of memory for array of `num` elements, each `size` bytes long
 - initializes memory to zero
- `void free(void* ptr):`
 - frees memory block, previously allocated by malloc, calloc, realloc, pointed by `ptr`
 - use exactly once for each pointer you allocate
- `size` argument:
 - *should* be computed using the `sizeof` operator
 - `sizeof`: takes a type and gives you its size
 - e.g., `sizeof(int), sizeof(int*)`

mem_mgmt.c

./mem_valgrind.sh

Memory Management Rules

- malloc what you free, free what you malloc
 - client should free memory allocated by client code
 - library should free memory allocated by library code
- Number mallocs = Number frees
 - Number mallocs > Number Frees: definitely a memory leak
 - Number mallocs < Number Frees: definitely a double free
- Free a malloc'ed block exactly once
 - Should not dereference a freed memory block
- Only malloc when necessary
 - Persistent, variable sized data structures
 - Concurrent accesses (we'll get there later in the semester)

Valgrind

- Find memory errors, detect memory leaks

- Common errors:

- Illegal read/write errors
 - Use of uninitialized values
 - Illegal frees
 - Overlapping source/destination addresses

- Typical solutions

- Did you allocate enough memory?
 - Did you accidentally free stack variables/something twice?
 - Did you initialize all your variables?
 - Did use something that you just free'd?

- leak-check=full

- Memcheck gives details for each definitely/possibly lost memory block (where it was allocated)



```

Terminal
File Edit View Terminal Tabs Help
[pwells2@newcell ~/junk]$ valgrind ./memleak
==16738== Memcheck, a memory error detector
==16738== Copyright (C) 2002-2010, and GNU GPL'd, by Julian Seward et al.
==16738== Using Valgrind-3.6.1 and LibVEX; rerun with -h for copyright info
==16738== Command: ./memleak
==16738==
==16738== Invalid write of size 4
==16738==    at 0x400589: main (mem_leak.c:32)
==16738==    Address 0x4c26068 is 0 bytes after a block of size 40 alloc'd
==16738==    at 0x4A0646F: malloc (vg_replace_malloc.c:236)
==16738==    by 0x400505: main (mem_leak.c:17)
==16738==
==16738== Invalid read of size 4
==16738==    at 0x400598: main (mem_leak.c:33)
==16738==    Address 0x4c26068 is 0 bytes after a block of size 40 alloc'd
==16738==    at 0x4A0646F: malloc (vg_replace_malloc.c:236)
==16738==    by 0x400505: main (mem_leak.c:17)
==16738==
==16738== HEAP SUMMARY:
==16738==    in use at exit: 410 bytes in 8 blocks
==16738==    total heap usage: 11 allocs, 3 frees, 590 bytes allocated
==16738==
==16738== LEAK SUMMARY:
==16738==    definitely lost: 410 bytes in 8 blocks
==16738==    indirectly lost: 0 bytes in 0 blocks
==16738==    possibly lost: 0 bytes in 0 blocks
==16738==    still reachable: 0 bytes in 0 blocks
==16738==    suppressed: 0 bytes in 0 blocks
==16738==
==16738== Rerun with --leak-check=full to see details of leaked memory
==16738==
==16738== For counts of detected and suppressed errors, rerun with: -v
==16738== ERROR SUMMARY: 36 errors from 2 contexts (suppressed: 4 from 4)
[pwells2@newcell ~/junk]$
  
```

Debugging

GDB

GDB

- No longer stepping through assembly!
Some GDB commands are different:
 - `si / si` → `step / next`
 - `break file.c:line_num`
 - `disas` → `list`
 - `print <any_var_name>` (in current frame)
- Use TUI mode (layout src)
 - Nice display for viewing source/executing commands
 - Buggy, so only use TUI mode to step through lines (no continue / finish)

Additional Topics

- Headers files and header guards
- Macros
- Appendix (C libraries)

Header Files

- Includes C declarations and macro definitions to be shared across multiple files
 - Only include function prototypes/macros; implementation code goes in .c file!
- Usage: `#include <header.h>`
 - `#include <lib>` for standard libraries (eg `#include <string.h>`)
 - `#include "file"` for your source files (eg `#include "header.h"`)
 - Never include .c files (bad practice)

```
// list.h
struct list_node {
    int data;
    struct list_node* next;
};
typedef struct list_node* node;
```

```
node new_list();
void add_node(int e, node l);
```

```
// list.c
#include "list.h"

node new_list() {
    // implementation
}
```

```
void add_node(int e, node l) {
    // implementation
}
```

```
// stacks.h
#include "list.h"
struct stack_head {
    node top;
    node bottom;
};
typedef struct stack_head* stack

stack new_stack();
void push(int e, stack S);
```

Header Guards

- Double-inclusion problem: include same header file twice

```
//grandfather.h
```

```
//father.h
```

```
#include "grandfather.h"
```

```
//child.h
```

```
#include "father.h"
```

```
#include "grandfather.h"
```

Error: child.h includes grandfather.h twice

- Solution: header guard ensures single inclusion

```
//grandfather.h
```

```
#ifndef GRANDFATHER_H
```

```
#define GRANDFATHER_H
```

```
#endif
```

```
//father.h
```

```
#ifndef FATHER_H
```

```
#define FATHER_H
```

```
#endif
```

```
//child.h
```

```
#include "father.h"
```

```
#include "grandfather.h"
```

Okay: child.h only includes grandfather.h once

Macros

./macros

- A way to replace a name with its macro definition
 - No function call overhead, type neutral
 - Think “find and replace” like in a text editor
- Uses:
 - defining constants (INT_MAX, ARRAY_SIZE)
 - defining simple operations (MAX(a, b))
 - 122-style contracts (REQUIRES, ENSURES)
- Warnings:
 - Use parentheses around arguments/expressions, to avoid problems after substitution
 - Do not pass expressions with side effects as arguments to macros

```
#define INT_MAX 0x7FFFFFFF
#define MAX(A, B) ((A) > (B) ? (A) : (B))
#define REQUIRES(COND) assert(COND)
#define WORD_SIZE 4
#define NEXT_WORD(a) ((char*)(a) + WORD_SIZE)
```


C Libraries

string.h: Common String/Array Methods

- One the most useful libraries available to you
- Used heavily in shell/proxy labs
- Important usage details regarding arguments:
 - prefixes: `str` -> strings, `mem` -> arbitrary memory blocks.
 - ensure that all strings are `'\0'` terminated!
 - ensure that `dest` is large enough to store `src`!
 - ensure that `src` actually contains `n` bytes!
 - ensure that `src/dest` don't overlap!



string.h: Common String/Array Methods

- Copying:

- `void *memcpy (void *dest, void *src, size_t n) :` copy n bytes of src into dest, return dest
- `char *strcpy(char *dest, char *src) :` copy src string into dest, return dest. *Make sure dest is large enough to contain src.*

- Concatenation:

- `char *strncat (char *dest, char *src, size_t n) :` append copy of src to end of dest reading at most n bytes, return dest
- `char *strcat (char *dest, char *src)` works for arbitrary length strings, but has the safety issues you've seen in attacklab

string.h: Common String/Array Methods (Continued)

- Comparison:

- `int strncmp (char *str1, char *str2, size_t n):` compare at most `n` bytes of `str1`, `str2` by character (based on ASCII value of each character, then string length), return comparison result
`str1 < str2: -1,`
`str1 == str2: 0,`
`str1 > str2: 1`
- `int strcmp(char *str1, char *str2):` compare `str1` to `str2`. *Make sure each string is long enough to be safely compared.*

string.h: Common String/Array Methods (Continued)

■ Searching:

- `char *strstr (char *str1, char *str2):` return pointer to *first* occurrence of `str2` in `str1`, else `NULL`
- `char *strtok (char *str, char *delimiters):` tokenize `str` according to delimiter characters provided in `delimiters`.
return the one token for each `strtok` call, using `str = NULL`

■ Other:

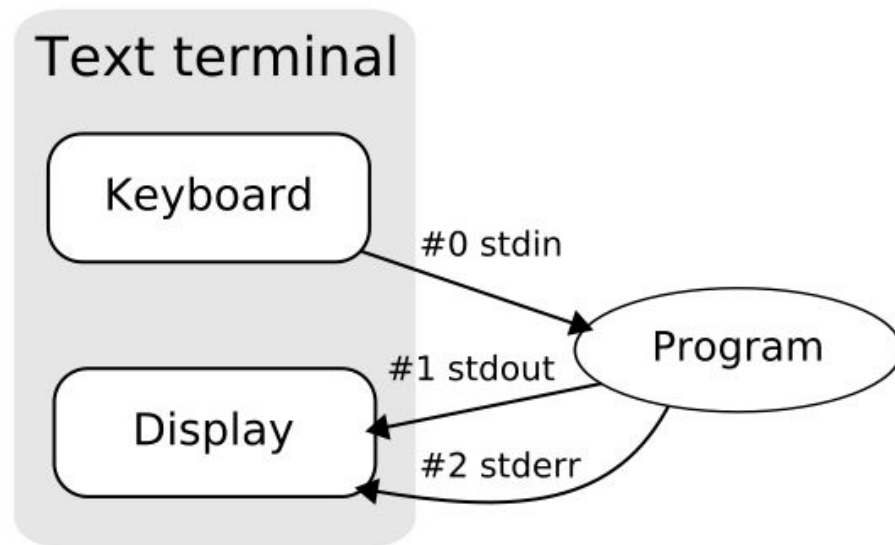
- `size_t strlen (const char *str):` returns length of the string (up to, but not including the `'\0'` character)
- `void *memset (void *ptr, int val, size_t n):` set first `n` bytes of memory block addressed by `ptr` to `val`
For *setting bytes only*. Don't use it to set or initialize `int` arrays, for example.

stdlib.h: General Purpose Functions

- Dynamic memory allocation:
 - `malloc`, `calloc`, `free`
- String conversion:
 - `int atoi(char *str)`: parse string into integral value (return 0 if not parsed)
- System Calls:
 - `void exit(int status)`: terminate calling process, return `status` to parent process
 - `void abort()`: aborts process abnormally
- Searching/Sorting:
 - provide array, array size, element size, comparator (function pointer)
 - `bsearch`: returns pointer to matching element in the array
 - `qsort`: sorts the array destructively
- Integer arithmetic:
 - `int abs(int n)`: returns absolute value of `n`
- Types:
 - `size_t`: unsigned integral type (store size of *any* object)

stdio.h

- Another really useful library.
- Used heavily in cache/shell/proxy labs
- Used for:
 - argument parsing
 - file handling
 - input/output
- `printf`, a fan favorite, comes from this library!



stdio.h: Common I/O Methods

- `FILE *fopen (char *filename, char *mode)`: open the file with specified filename in specified mode (read, write, append, etc), associate it with stream identified by returned file pointer
- `int fscanf (FILE *stream, char *format, ...)`: read data from the stream, store it according to the parameter format at the memory locations pointed at by additional arguments.
- `int fclose (FILE *stream)`: close the file associated with stream
- `int fprintf (FILE *stream, char *format, ...)`: write the C string pointed at by format to the stream, using any additional arguments to fill in format specifiers.
- `fgets`

Getopt

- Need to include `unistd.h` to use
- Used to parse command-line arguments.
- Typically called in a loop to retrieve arguments
- Switch statement used to handle options
 - colon indicates required argument
 - `optarg` is set to value of option argument
- Returns -1 when no more arguments present
- See recitation 6 slides for more examples

```
int main(int argc, char **argv)
{
    int opt, x;
    /* looping over arguments */
    while((opt=getopt(argc,argv,"x:"))>0){
        switch(opt) {
            case 'x':
                x = atoi(optarg);
                break;
            default:
                printf("wrong argument\n");
                break;
        }
    }
}
```

Note about Library Functions

- These functions can return error codes
 - `malloc` could fail
 - ```
int x;
if ((x = malloc(sizeof(int))) == NULL)
 printf("Malloc failed!!!\n");
```
  - a file couldn't be opened
  - a string may be incorrectly parsed
- Remember to check for the error cases and handle the errors accordingly
  - may have to terminate the program (eg `malloc` fails)
  - may be able to recover (user entered bad input)

# Style

- Documentation
  - file header, function header, comments
- Variable Names & Magic Numbers
  - `new_cache_size` is good, not `new_cacheSize` or `size`
  - Use `#define CACHESIZE 128`
- Modularity
  - helper functions
- Error Checking
  - `malloc`, library functions...
- Memory & File Handling
  - free memory, close files
- Check [style guide](#) for detailed information